

Chapter 13. Data Analysis Examples

Now that we've reached the final chapter of this book, we're going to take a look at a number of real-world datasets. For each dataset, we'll use the techniques presented in this book to extract meaning from the raw data. The demonstrated techniques can be applied to all manner of other datasets. This chapter contains a collection of miscellaneous example datasets that you can use for practice with the tools in this book.

The example datasets are found in the book's accompanying [GitHub repository](#). If you are unable to access GitHub, you can also get them from the [repository mirror on Gitee](#).

13.1 Bitly Data from 1.USA.gov

In 2011, the URL shortening service [Bitly](#) partnered with the US government website [USA.gov](#) to provide a feed of anonymous data gathered from users who shorten links ending with `.gov` or `.mil`. In 2011, a live feed as well as hourly snapshots were available as downloadable text files. This service is shut down at the time of this writing (2022), but we preserved one of the data files for the book's examples.

In the case of the hourly snapshots, each line in each file contains a common form of web data known as JSON, which stands for JavaScript Object Notation. For example, if we read just the first line of a file, we may see something like this:

```
In [5]: path = "datasets/bitly_usagov/example.txt"

In [6]: with open(path) as f:
...:     print(f.readline())
...:

{ "a": "Mozilla\\5.0 (Windows NT 6.1; WOW64) AppleWebKit\\535.11
(KHTML, like Gecko) Chrome\\17.0.963.78 Safari\\535.11", "c": "US", "nk": "1
"tz": "America\\New_York", "gr": "MA", "g": "A6qOVH", "h": "wfLQtf", "l": "o
"orofrog", "al": "en-US,en;q=0.8", "hh": "1.usa.gov", "r": "http:\\\\www.facebook.com\\l\\7AQEFzjSi\\1.usa.gov\\wfLQtf", "u": "http:\\\\www.ncbi.nlm.nih.gov\\pubmed\\22415991", "t": 1331923247, "hc": 1331822918, "cy": "Danvers", "ll": [ 42.576698, -70.954903 ] }
```

Python has both built-in and third-party libraries for converting a JSON string into a Python dictionary. Here we'll use the `json` module and its `loads` function invoked on each line in the sample file we downloaded:

```
import json
with open(path) as f:
    records = [json.loads(line) for line in f]
```

The resulting object `records` is now a list of Python dictionaries:

```
In [18]: records[0]
Out[18]:
{'a': 'Mozilla/5.0 (Windows NT 6.1; WOW64) AppleWebKit/535.11 (KHTML, like Gecko; Chrome/17.0.963.78 Safari/535.11',
 'al': 'en-US,en;q=0.8',
 'c': 'US',
 'cy': 'Danvers',
 'g': 'A6q0VH',
 'gr': 'MA',
 'h': 'wfLQtf',
 'hc': 1331822918,
 'hh': '1.usa.gov',
 'l': 'orofrog',
 'll': [42.576698, -70.954903],
 'nk': 1,
 'r': 'http://www.facebook.com/l/7AQEFzjSi/1.usa.gov/wfLQtf',
 't': 1331923247,
 'tz': 'America/New_York',
 'u': 'http://www.ncbi.nlm.nih.gov/pubmed/22415991'}
```

Counting Time Zones in Pure Python

Suppose we were interested in finding the time zones that occur most often in the dataset (the `tz` field). There are many ways we could do this. First, let's extract a list of time zones again using a list comprehension:

```
In [15]: time_zones = [rec["tz"] for rec in records]
```

```
-----
KeyError                                     Traceback (most recent call last)
<ipython-input-15-abdeba901c13> in <module>
      1 time_zones = [rec["tz"] for rec in records]
<ipython-input-15-abdeba901c13> in <listcomp>(.0)
      1 time_zones = [rec["tz"] for rec in records]
KeyError: 'tz'
```

Oops! Turns out that not all of the records have a time zone field. We can handle this by adding the check `if "tz" in rec` at the end of the list comprehension:

```
In [16]: time_zones = [rec["tz"] for rec in records if "tz" in rec]

In [17]: time_zones[:10]
Out[17]:
['America/New_York',
 'America/Denver',
 'America/New_York',
 'America/Sao_Paulo',
 'America/New_York',
 'America/New_York',
 'Europe/Warsaw',
 '',
 '',
 '']
```

Just looking at the first 10 time zones, we see that some of them are unknown (empty string). You can filter these out also, but I'll leave them in for now. Next, to produce counts by time zone, I'll show two approaches: a harder way (using just the Python standard library) and a simpler way (using pandas). One way to do the counting is to use a dictionary to store counts while we iterate through the time zones:

```
def get_counts(sequence):
    counts = {}
    for x in sequence:
        if x in counts:
            counts[x] += 1
        else:
            counts[x] = 1
    return counts
```

Using more advanced tools in the Python standard library, you can write the same thing more briefly:

```
from collections import defaultdict

def get_counts2(sequence):
    counts = defaultdict(int) # values will initialize to 0
    for x in sequence:
```

```
counts[x] += 1
return counts
```

I put this logic in a function just to make it more reusable. To use it on the time zones, just pass the `time_zones` list:

```
In [20]: counts = get_counts(time_zones)

In [21]: counts["America/New_York"]
Out[21]: 1251

In [22]: len(time_zones)
Out[22]: 3440
```

If we wanted the top 10 time zones and their counts, we can make a list of tuples by `(count, timezone)` and sort it:

```
def top_counts(count_dict, n=10):
    value_key_pairs = [(count, tz) for tz, count in count_dict.items()]
    value_key_pairs.sort()
    return value_key_pairs[-n:]
```

We have then:

```
In [24]: top_counts(counts)
Out[24]:
[(33, 'America/Sao_Paulo'),
 (35, 'Europe/Madrid'),
 (36, 'Pacific/Honolulu'),
 (37, 'Asia/Tokyo'),
 (74, 'Europe/London'),
 (191, 'America/Denver'),
 (382, 'America/Los_Angeles'),
 (400, 'America/Chicago'),
 (521, ''),
 (1251, 'America/New_York')]
```

If you search the Python standard library, you may find the `collections.Counter` class, which makes this task even simpler:

```
In [25]: from collections import Counter

In [26]: counts = Counter(time_zones)
```

```
In [27]: counts.most_common(10)
Out[27]:
[('America/New_York', 1251),
 ('', 521),
 ('America/Chicago', 400),
 ('America/Los_Angeles', 382),
 ('America/Denver', 191),
 ('Europe/London', 74),
 ('Asia/Tokyo', 37),
 ('Pacific/Honolulu', 36),
 ('Europe/Madrid', 35),
 ('America/Sao_Paulo', 33)]
```

Counting Time Zones with pandas

You can create a DataFrame from the original set of records by passing the list of records to `pandas.DataFrame`:

```
In [28]: frame = pd.DataFrame(records)
```

We can look at some basic information about this new DataFrame, such as column names, inferred column types, or number of missing values, using `frame.info()`:

```
In [29]: frame.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3560 entries, 0 to 3559
Data columns (total 18 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   a            3440 non-null    object 
 1   c            2919 non-null    object 
 2   nk           3440 non-null    float64
 3   tz           3440 non-null    object 
 4   gr           2919 non-null    object 
 5   g            3440 non-null    object 
 6   h            3440 non-null    object 
 7   l            3440 non-null    object 
 8   al           3094 non-null    object 
 9   hh           3440 non-null    object 
 10  r            3440 non-null    object 
 11  u            3440 non-null    object 
 12  t            3440 non-null    float64
 13  hc           3440 non-null    float64
```

```
14 cy 2919 non-null object
15 ll 2919 non-null object
16 _heartbeat_ 120 non-null float64
17 kw 93 non-null object
dtypes: float64(4), object(14)
memory usage: 500.8+ KB
```

```
In [30]: frame["tz"].head()
Out[30]:
0 America/New_York
1 America/Denver
2 America/New_York
3 America/Sao_Paulo
4 America/New_York
Name: tz, dtype: object
```

The output shown for the `frame` is the *summary view*, shown for large DataFrame objects. We can then use the `value_counts` method for the Series:

```
In [31]: tz_counts = frame["tz"].value_counts()

In [32]: tz_counts.head()
Out[32]:
America/New_York      1251
                      521
America/Chicago        400
America/Los_Angeles   382
America/Denver         191
Name: tz, dtype: int64
```

We can visualize this data using matplotlib. We can make the plots a bit nicer by filling in a substitute value for unknown or missing time zone data in the records. We replace the missing values with the `fillna` method and use Boolean array indexing for the empty strings:

```
In [33]: clean_tz = frame["tz"].fillna("Missing")

In [34]: clean_tz[clean_tz == ""] = "Unknown"

In [35]: tz_counts = clean_tz.value_counts()

In [36]: tz_counts.head()
Out[36]:
America/New_York      1251
```

```
Unknown          521
America/Chicago   400
America/Los_Angeles 382
America/Denver     191
Name: tz, dtype: int64
```

At this point, we can use the [seaborn package](#) to make a horizontal bar plot (see [Figure 13-1](#) for the resulting visualization):

```
In [38]: import seaborn as sns

In [39]: subset = tz_counts.head()

In [40]: sns.barplot(y=subset.index, x=subset.to_numpy())
```

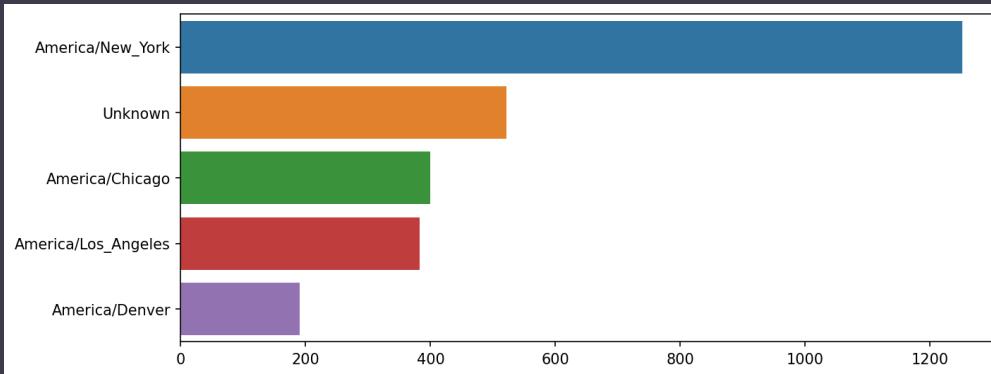


Figure 13-1. Top time zones in the 1.usa.gov sample data

The `a` field contains information about the browser, device, or application used to perform the URL shortening:

```
In [41]: frame["a"][1]
Out[41]: 'GoogleMaps/RochesterNY'

In [42]: frame["a"][50]
Out[42]: 'Mozilla/5.0 (Windows NT 5.1; rv:10.0.2) Gecko/20100101 Firefox/10.0.2'

In [43]: frame["a"][51][:50] # long line
Out[43]: 'Mozilla/5.0 (Linux; U; Android 2.2.2; en-us; LG-P99)'
```

Parsing all of the interesting information in these “agent” strings may seem like a daunting task. One possible strategy is to split off the first token in the string (corresponding roughly to the browser capability) and make another summary of the user behavior:

```
In [44]: results = pd.Series([x.split()[0] for x in frame["a"].dropna()])

In [45]: results.head(5)
Out[45]:
0           Mozilla/5.0
1    GoogleMaps/RochesterNY
2           Mozilla/4.0
3           Mozilla/5.0
4           Mozilla/5.0
dtype: object

In [46]: results.value_counts().head(8)
Out[46]:
Mozilla/5.0            2594
Mozilla/4.0              601
GoogleMaps/RochesterNY      121
Opera/9.80                  34
TEST_INTERNET_AGENT          24
GoogleProducer                  21
Mozilla/6.0                      5
BlackBerry8520/5.0.0.681          4
dtype: int64
```

Now, suppose you wanted to decompose the top time zones into Windows and non-Windows users. As a simplification, let's say that a user is on Windows if the string "Windows" is in the agent string. Since some of the agents are missing, we'll exclude these from the data:

```
In [47]: cframe = frame[frame["a"].notna()].copy()
```

We want to then compute a value for whether or not each row is Windows:

```
In [48]: cframe["os"] = np.where(cframe["a"].str.contains("Windows"),
...:                         "Windows", "Not Windows")

In [49]: cframe["os"].head(5)
Out[49]:
0       Windows
1    Not Windows
2       Windows
3    Not Windows
4       Windows
Name: os, dtype: object
```

Then, you can group the data by its time zone column and this new list of operating systems:

```
In [50]: by_tz_os = cframe.groupby(["tz", "os"])
```

The group counts, analogous to the `value_counts` function, can be computed with `size`. This result is then reshaped into a table with `unstack`:

```
In [51]: agg_counts = by_tz_os.size().unstack().fillna(0)
```

```
In [52]: agg_counts.head()
```

```
Out[52]:
```

os	Not	Windows	Windows
tz		245.0	276.0
Africa/Cairo	0.0	3.0	
Africa/Casablanca	0.0	1.0	
Africa/Ceuta	0.0	2.0	
Africa/Johannesburg	0.0	1.0	

Finally, let's select the top overall time zones. To do so, I construct an indirect index array from the row counts in `agg_counts`. After computing the row counts with `agg_counts.sum("columns")`, I can call `argsort()` to obtain an index array that can be used to sort in ascending order:

```
In [53]: indexer = agg_counts.sum("columns").argsort()
```

```
In [54]: indexer.values[:10]
```

```
Out[54]: array([24, 20, 21, 92, 87, 53, 54, 57, 26, 55])
```

I use `take` to select the rows in that order, then slice off the last 10 rows (largest values):

```
In [55]: count_subset = agg_counts.take(indexer[-10:])
```

```
In [56]: count_subset
```

```
Out[56]:
```

os	Not	Windows	Windows
tz			
America/Sao_Paulo	13.0	20.0	
Europe/Madrid	16.0	19.0	

Pacific/Honolulu	0.0	36.0
Asia/Tokyo	2.0	35.0
Europe/London	43.0	31.0
America/Denver	132.0	59.0
America/Los_Angeles	130.0	252.0
America/Chicago	115.0	285.0
	245.0	276.0
America/New_York	339.0	912.0

pandas has a convenience method called `nlargest` that does the same thing:

```
In [57]: agg_counts.sum(axis="columns").nlargest(10)
Out[57]:
tz
America/New_York      1251.0
                      521.0
America/Chicago        400.0
America/Los_Angeles   382.0
America/Denver         191.0
Europe/London          74.0
Asia/Tokyo             37.0
Pacific/Honolulu       36.0
Europe/Madrid           35.0
America/Sao_Paulo      33.0
dtype: float64
```

Then, this can be plotted in a grouped bar plot comparing the number of Windows and non-Windows users, using seaborn's `barplot` function (see [Figure 13-2](#)). I first call `count_subset.stack()` and reset the index to rearrange the data for better compatibility with seaborn:

```
In [59]: count_subset = count_subset.stack()

In [60]: count_subset.name = "total"

In [61]: count_subset = count_subset.reset_index()

In [62]: count_subset.head(10)
Out[62]:
            tz          os  total
0  America/Sao_Paulo  Not Windows   13.0
1  America/Sao_Paulo      Windows   20.0
2    Europe/Madrid  Not Windows   16.0
3    Europe/Madrid      Windows   19.0
```

```

4   Pacific/Honolulu  Not Windows    0.0
5   Pacific/Honolulu      Windows   36.0
6       Asia/Tokyo     Not Windows    2.0
7       Asia/Tokyo      Windows   35.0
8   Europe/London     Not Windows   43.0
9   Europe/London      Windows   31.0

```

```
In [63]: sns.barplot(x="total", y="tz", hue="os", data=count_subset)
```

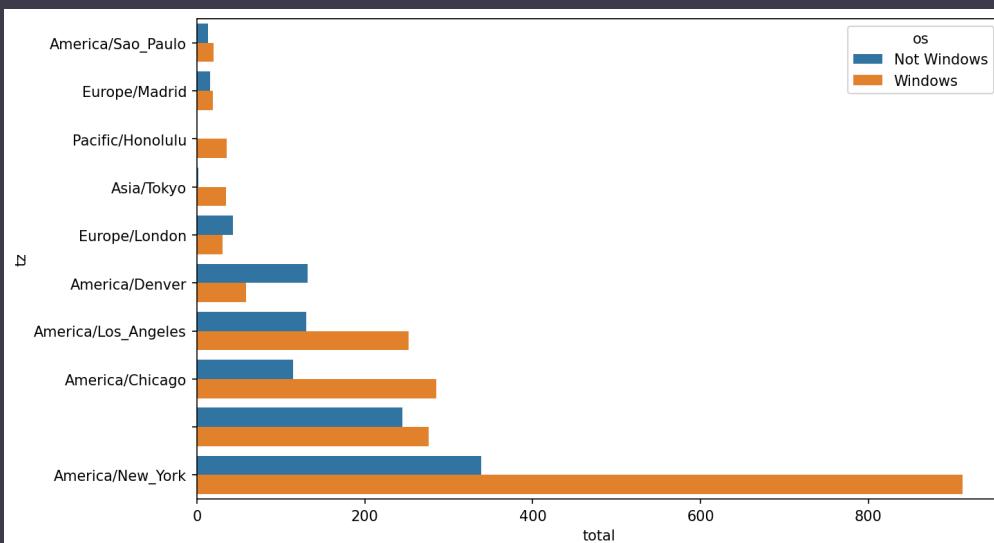


Figure 13-2. Top time zones by Windows and non-Windows users

It is a bit difficult to see the relative percentage of Windows users in the smaller groups, so let's normalize the group percentages to sum to 1:

```

def norm_total(group):
    group["normed_total"] = group["total"] / group["total"].sum()
    return group

results = count_subset.groupby("tz").apply(norm_total)

```

Then plot this in [Figure 13-3](#):

```
In [66]: sns.barplot(x="normed_total", y="tz", hue="os", data=results)
```

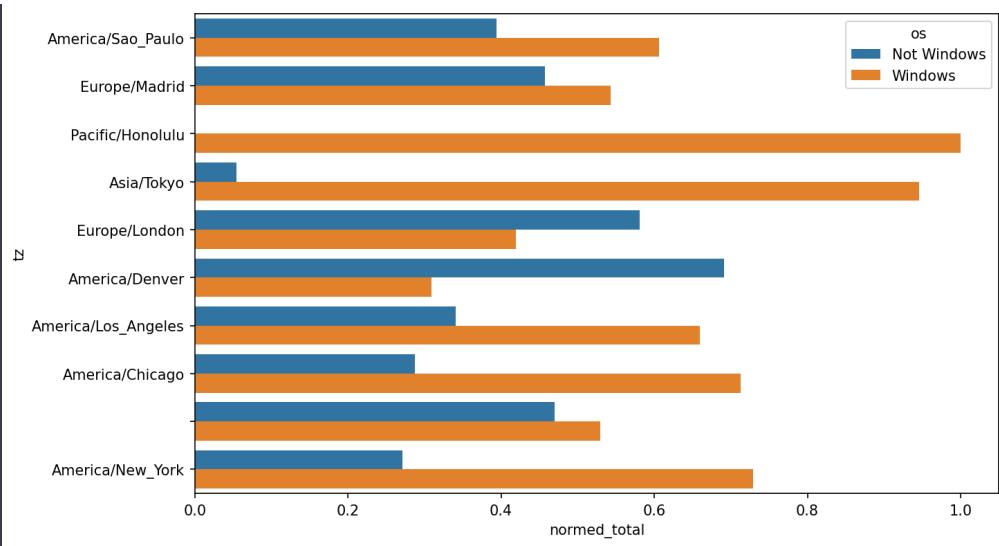


Figure 13-3. Percentage Windows and non-Windows users in top occurring time zones

We could have computed the normalized sum more efficiently by using the `transform` method with `groupby`:

```
In [67]: g = count_subset.groupby("tz")
In [68]: results2 = count_subset["total"] / g["total"].transform("sum")
```

13.2 MovieLens 1M Dataset

[GroupLens Research](#) provides a number of collections of movie ratings data collected from users of MovieLens in the late 1990s and early 2000s. The data provides movie ratings, movie metadata (genres and year), and demographic data about the users (age, zip code, gender identification, and occupation). Such data is often of interest in the development of recommendation systems based on machine learning algorithms. While we do not explore machine learning techniques in detail in this book, I will show you how to slice and dice datasets like these into the exact form you need.

The MovieLens 1M dataset contains one million ratings collected from six thousand users on four thousand movies. It's spread across three tables: ratings, user information, and movie information. We can load each table into a pandas DataFrame object using `pandas.read_table`. Run the following code in a Jupyter cell:

```
unames = ["user_id", "gender", "age", "occupation", "zip"]
users = pd.read_table("datasets/movielens/users.dat", sep="::",
```

```

header=None, names=unames, engine="python"))

rnames = ["user_id", "movie_id", "rating", "timestamp"]
ratings = pd.read_table("datasets/movielens/ratings.dat", sep="::",
                        header=None, names=rnames, engine="python"))

mnames = ["movie_id", "title", "genres"]
movies = pd.read_table("datasets/movielens/movies.dat", sep="::",
                        header=None, names=mnames, engine="python"))

```

You can verify that everything succeeded by looking at each DataFrame:

```

In [70]: users.head(5)
Out[70]:
   user_id  gender  age occupation      zip
0         1        F    1          10  48067
1         2        M   56          16  70072
2         3        M   25          15  55117
3         4        M   45           7  02460
4         5        M   25          20  55455

In [71]: ratings.head(5)
Out[71]:
   user_id  movie_id  rating  timestamp
0         1       1193      5  978300760
1         1        661      3  978302109
2         1        914      3  978301968
3         1       3408      4  978300275
4         1       2355      5  978824291

In [72]: movies.head(5)
Out[72]:
   movie_id                      title   genre
0         1             Toy Story (1995) Animation|Children's|Comedy
1         2             Jumanji (1995) Adventure|Children's|Fantasy
2         3  Grumpier Old Men (1995) Comedy|Romance
3         4        Waiting to Exhale (1995) Comedy|Drama
4         5 Father of the Bride Part II (1995) Comedy

In [73]: ratings
Out[73]:
   user_id  movie_id  rating  timestamp
0         1       1193      5  978300760
1         1        661      3  978302109
2         1        914      3  978301968
3         1       3408      4  978300275
4         1       2355      5  978824291

```

```

1000204      6040      1091      1  956716541
1000205      6040      1094      5  956704887
1000206      6040      562       5  956704746
1000207      6040      1096      4  956715648
1000208      6040      1097      4  956715569
[1000209 rows x 4 columns]

```

Note that ages and occupations are coded as integers indicating groups described in the dataset's *README* file. Analyzing the data spread across three tables is not a simple task; for example, suppose you wanted to compute mean ratings for a particular movie by gender identity and age. As you will see, this is more convenient to do with all of the data merged together into a single table. Using pandas's `merge` function, we first merge `ratings` with `users` and then merge that result with the `movies` data. pandas infers which columns to use as the merge (or *join*) keys based on overlapping names:

```
In [74]: data = pd.merge(pd.merge(ratings, users), movies)
```

In [75]: data

Out[75]:

	user_id	movie_id	rating	timestamp	gender	age	occupation	zip
0		1	1193	5	978300760	F	1	10 48067
1		2	1193	5	978298413	M	56	16 70072
2		12	1193	4	978220179	M	25	12 32793
3		15	1193	4	978199279	M	25	7 22903
4		17	1193	5	978158471	M	50	1 95350
	
1000204	5949	2198	5	958846401	M	18	17 47901	
1000205	5675	2703	3	976029116	M	35	14 30030	
1000206	5780	2845	1	958153068	M	18	17 92886	
1000207	5851	3607	5	957756608	F	18	20 55410	
1000208	5938	2909	4	957273353	M	25	1 35401	
					title		genres	
0		One Flew Over the Cuckoo's Nest (1975)					Drama	
1		One Flew Over the Cuckoo's Nest (1975)					Drama	
2		One Flew Over the Cuckoo's Nest (1975)					Drama	
3		One Flew Over the Cuckoo's Nest (1975)					Drama	
4		One Flew Over the Cuckoo's Nest (1975)					Drama	
		
1000204			Modulations (1998)				Documentary	
1000205			Broken Vessels (1998)				Drama	
1000206			White Boys (1999)				Drama	
1000207			One Little Indian (1973)			Comedy	Drama	Western
1000208		Five Wives, Three Secretaries and Me (1998)					Documentary	

```
[1000209 rows x 10 columns]
```

```
In [76]: data.iloc[0]
Out[76]:
user_id                               1
movie_id                            1193
rating                                5
timestamp                         978300760
gender                                 F
age                                     1
occupation                           10
zip                                    48067
title      One Flew Over the Cuckoo's Nest (1975)
genres                                Drama
Name: 0, dtype: object
```

To get mean movie ratings for each film grouped by gender, we can use the `pivot_table` method:

```
In [77]: mean_ratings = data.pivot_table("rating", index="title",
                                         ....:                                         columns="gender", aggfunc="mean")

In [78]: mean_ratings.head(5)
Out[78]:
gender          F          M
title
$1,000,000 Duck (1971)    3.375000  2.761905
'Night Mother (1986)       3.388889  3.352941
'Til There Was You (1997)  2.675676  2.733333
'burbs, The (1989)        2.793478  2.962085
...And Justice for All (1979) 3.828571  3.689024
```

This produced another DataFrame containing mean ratings with movie titles as row labels (the “index”) and gender as column labels. I first filter down to movies that received at least 250 ratings (an arbitrary number); to do this, I group the data by title, and use `size()` to get a Series of group sizes for each title:

```
In [79]: ratings_by_title = data.groupby("title").size()

In [80]: ratings_by_title.head()
Out[80]:
title
$1,000,000 Duck (1971)      37
'Night Mother (1986)         70
```

```
'Til There Was You (1997)      52
'burbs, The (1989)            303
...And Justice for All (1979)  199
dtype: int64
```

```
In [81]: active_titles = ratings_by_title.index[ratings_by_title >= 250]

In [82]: active_titles
Out[82]:
Index([''burbs, The (1989)', '10 Things I Hate About You (1999)',
       '101 Dalmatians (1961)', '101 Dalmatians (1996)', '12 Angry Men (1957',
       '13th Warrior, The (1999)', '2 Days in the Valley (1996)', '20,000 Leagues Under the Sea (1954)', '2001: A Space Odyssey (1968)',
       '2010 (1984)', ...
       'X-Men (2000)', 'Year of Living Dangerously (1982)', 'Yellow Submarine (1968)', 'You've Got Mail (1998)', 'Young Frankenstein (1974)', 'Young Guns (1988)', 'Young Guns II (1990)', 'Young Sherlock Holmes (1985)', 'Zero Effect (1998)', 'eXistenZ (1999)'],
      dtype='object', name='title', length=1216)
```

The index of titles receiving at least 250 ratings can then be used to select rows from `mean_ratings` using `.loc`:

```
In [83]: mean_ratings = mean_ratings.loc[active_titles]

In [84]: mean_ratings
Out[84]:
gender          F          M
title
'burbs, The (1989)      2.793478  2.962085
10 Things I Hate About You (1999)  3.646552  3.311966
101 Dalmatians (1961)     3.791444  3.500000
101 Dalmatians (1996)     3.240000  2.911215
12 Angry Men (1957)       4.184397  4.328421
...
Young Guns (1988)        3.371795  3.425620
Young Guns II (1990)      2.934783  2.904025
Young Sherlock Holmes (1985) 3.514706  3.363344
Zero Effect (1998)        3.864407  3.723140
eXistenZ (1999)          3.098592  3.289086
[1216 rows x 2 columns]
```

To see the top films among female viewers, we can sort by the `F` column in descending order:

```
In [86]: top_female_ratings = mean_ratings.sort_values("F", ascending=False)

In [87]: top_female_ratings.head()
Out[87]:
gender          F          M
title
Close Shave, A (1995)  4.644444  4.473795
Wrong Trousers, The (1993)  4.588235  4.478261
Sunset Blvd. (a.k.a. Sunset Boulevard) (1950)  4.572650  4.464589
Wallace & Gromit: The Best of Aardman Animation (1996)  4.563107  4.385075
Schindler's List (1993)  4.562602  4.491415
```

Measuring Rating Disagreement

Suppose you wanted to find the movies that are most divisive between male and female viewers. One way is to add a column to `mean_ratings` containing the difference in means, then sort by that:

```
In [88]: mean_ratings["diff"] = mean_ratings["M"] - mean_ratings["F"]
```

Sorting by `"diff"` yields the movies with the greatest rating difference so that we can see which ones were preferred by women:

```
In [89]: sorted_by_diff = mean_ratings.sort_values("diff")

In [90]: sorted_by_diff.head()
Out[90]:
gender          F          M      diff
title
Dirty Dancing (1987)  3.790378  2.959596 -0.830782
Jumpin' Jack Flash (1986)  3.254717  2.578358 -0.676359
Grease (1978)  3.975265  3.367041 -0.608224
Little Women (1994)  3.870588  3.321739 -0.548849
Steel Magnolias (1989)  3.901734  3.365957 -0.535777
```

Reversing the order of the rows and again slicing off the top 10 rows, we get the movies preferred by men that women didn't rate as highly:

```
In [91]: sorted_by_diff[::-1].head()
Out[91]:
gender          F          M      diff
title
Good, The Bad and The Ugly, The (1966)  3.494949  4.221300  0.726351
```

Kentucky Fried Movie, The (1977)	2.878788	3.555147	0.676359
Dumb & Dumber (1994)	2.697987	3.336595	0.638608
Longest Day, The (1962)	3.411765	4.031447	0.619682
Cable Guy, The (1996)	2.250000	2.863787	0.613787

Suppose instead you wanted the movies that elicited the most disagreement among viewers, independent of gender identification. Disagreement can be measured by the variance or standard deviation of the ratings. To get this, we first compute the rating standard deviation by title and then filter down to the active titles:

```
In [92]: rating_std_by_title = data.groupby("title")["rating"].std()

In [93]: rating_std_by_title = rating_std_by_title.loc[active_titles]

In [94]: rating_std_by_title.head()
Out[94]:
title
'burbs, The (1989)           1.107760
10 Things I Hate About You (1999) 0.989815
101 Dalmatians (1961)        0.982103
101 Dalmatians (1996)        1.098717
12 Angry Men (1957)          0.812731
Name: rating, dtype: float64
```

Then, we sort in descending order and select the first 10 rows, which are roughly the 10 most divisively rated movies:

```
In [95]: rating_std_by_title.sort_values(ascending=False)[:10]
Out[95]:
title
Dumb & Dumber (1994)           1.321333
Blair Witch Project, The (1999) 1.316368
Natural Born Killers (1994)     1.307198
Tank Girl (1995)                1.277695
Rocky Horror Picture Show, The (1975) 1.260177
Eyes Wide Shut (1999)             1.259624
Evita (1996)                     1.253631
Billy Madison (1995)              1.249970
Fear and Loathing in Las Vegas (1998) 1.246408
Bicentennial Man (1999)            1.245533
Name: rating, dtype: float64
```

You may have noticed that movie genres are given as a pipe-separated (1) string, since a single movie can belong to multiple genres. To help us group the ratings data by genre, we can use the `explode` method on DataFrame. Let's take a look at how this works. First, we can split the genres string into a list of genres using the `str.split` method on the Series:

```
In [96]: movies["genres"].head()
Out[96]:
0      Animation|Children's|Comedy
1      Adventure|Children's|Fantasy
2                  Comedy|Romance
3                  Comedy|Drama
4                  Comedy
Name: genres, dtype: object

In [97]: movies["genres"].head().str.split("|")
Out[97]:
0      [Animation, Children's, Comedy]
1      [Adventure, Children's, Fantasy]
2      [Comedy, Romance]
3      [Comedy, Drama]
4      [Comedy]
Name: genres, dtype: object

In [98]: movies["genre"] = movies.pop("genres").str.split("|")

In [99]: movies.head()
Out[99]:
   movie_id          title \
0        1      Toy Story (1995)
1        2           Jumanji (1995)
2        3  Grumpier Old Men (1995)
3        4      Waiting to Exhale (1995)
4        5  Father of the Bride Part II (1995)

                  genre
0      [Animation, Children's, Comedy]
1      [Adventure, Children's, Fantasy]
2      [Comedy, Romance]
3      [Comedy, Drama]
4      [Comedy]
```

Now, calling `movies.explode("genre")` generates a new DataFrame with one row for each “inner” element in each list of movie genres. For example, if a movie is classified as both a comedy and a romance, then there will be two rows in the result, one with just “Comedy” and the other with just “Romance”:

```
In [100]: movies_exploded = movies.explode("genre")

In [101]: movies_exploded[:10]
Out[101]:
   movie_id          title      genre
0         1    Toy Story (1995)  Animation
0         1    Toy Story (1995) Children's
0         1    Toy Story (1995)    Comedy
1         2       Jumanji (1995) Adventure
1         2       Jumanji (1995) Children's
1         2       Jumanji (1995)    Fantasy
2         3  Grumpier Old Men (1995)    Comedy
2         3  Grumpier Old Men (1995) Romance
3         4  Waiting to Exhale (1995)    Comedy
3         4  Waiting to Exhale (1995)    Drama
```

Now, we can merge all three tables together and group by genre:

```
In [102]: ratings_with_genre = pd.merge(pd.merge(movies_exploded, ratings),
                                         )

In [103]: ratings_with_genre.iloc[0]
Out[103]:
   movie_id          1
   title       Toy Story (1995)
   genre        Animation
   user_id          1
   rating          5
   timestamp     978824268
   gender          F
   age             1
   occupation      10
   zip            48067
   Name: 0, dtype: object

In [104]: genre_ratings = (ratings_with_genre.groupby(["genre", "age"])
                           .mean()
                           .unstack("age"))

In [105]: genre_ratings[:10]
Out[105]:
   age           1          18          25          35          45          50
   genre
   Action      3.506385  3.447097  3.453358  3.538107  3.528543  3.611333
   Adventure   3.449975  3.408525  3.443163  3.515291  3.528963  3.628163
   Animation   3.476113  3.624014  3.701228  3.740545  3.734856  3.780020
   Children's  3.241642  3.294257  3.426873  3.518423  3.527593  3.556555
```

Comedy	3.497491	3.460417	3.490385	3.561984	3.591789	3.646868
Crime	3.710170	3.668054	3.680321	3.733736	3.750661	3.810688
Documentary	3.730769	3.865865	3.946690	3.953747	3.966521	3.908108
Drama	3.794735	3.721930	3.726428	3.782512	3.784356	3.878415
Fantasy	3.317647	3.353778	3.452484	3.482301	3.532468	3.581570
Film-Noir	4.145455	3.997368	4.058725	4.064910	4.105376	4.175401
age		56				
genre						
Action	3.610709					
Adventure	3.649064					
Animation	3.756233					
Children's	3.621822					
Comedy	3.650949					
Crime	3.832549					
Documentary	3.961538					
Drama	3.933465					
Fantasy	3.532700					
Film-Noir	4.125932					

13.3 US Baby Names 1880–2010

The United States Social Security Administration (SSA) has made available data on the frequency of baby names from 1880 through the present. Hadley Wickham, an author of several popular R packages, has this dataset in illustrating data manipulation in R.

We need to do some data wrangling to load this dataset, but once we do that we will have a DataFrame that looks like this:

```
In [4]: names.head(10)
Out[4]:
      name  sex  births   year
0     Mary    F     7065  1880
1     Anna    F     2604  1880
2     Emma    F     2003  1880
3  Elizabeth    F     1939  1880
4    Minnie    F     1746  1880
5  Margaret    F     1578  1880
6      Ida    F     1472  1880
7     Alice    F     1414  1880
8    Bertha    F     1320  1880
9     Sarah    F     1288  1880
```

There are many things you might want to do with the dataset:

- Visualize the proportion of babies given a particular name (your own, or another name) over time
- Determine the relative rank of a name
- Determine the most popular names in each year or the names whose popularity has advanced or declined the most
- Analyze trends in names: vowels, consonants, length, overall diversity, changes in spelling, first and last letters
- Analyze external sources of trends: biblical names, celebrities, demographics

With the tools in this book, many of these kinds of analyses are within reach, so I will walk you through some of them.

As of this writing, the US Social Security Administration makes available data files, one per year, containing the total number of births for each sex/name combination. You can download the [raw archive](#) of these files.

If this page has been moved by the time you’re reading this, it can most likely be located again with an internet search. After downloading the “National data” file *names.zip* and unzipping it, you will have a directory containing a series of files like *yob1880.txt*. I use the Unix `head` command to look at the first 10 lines of one of the files (on Windows, you can use the `more` command or open it in a text editor):

```
In [106]: !head -n 10 datasets/babynames/yob1880.txt
Mary,F,7065
Anna,F,2604
Emma,F,2003
Elizabeth,F,1939
Minnie,F,1746
Margaret,F,1578
Ida,F,1472
Alice,F,1414
Bertha,F,1320
Sarah,F,1288
```

As this is already in comma-separated form, it can be loaded into a DataFrame with `pandas.read_csv`:

```
In [107]: names1880 = pd.read_csv("datasets/babynames/yob1880.txt",
... . . . .
In [108]: names1880
```

```
Out[108]:
```

		name	sex	births
0		Mary	F	7065
1		Anna	F	2604
2		Emma	F	2003
3		Elizabeth	F	1939
4		Minnie	F	1746
...	
1995		Woodie	M	5
1996		Worthy	M	5
1997		Wright	M	5
1998		York	M	5
1999		Zachariah	M	5
[2000	rows	x	3	columns]

These files only contain names with at least five occurrences in each year, so for simplicity's sake we can use the sum of the births column by sex as the total number of births in that year:

```
In [109]: names1880.groupby("sex")["births"].sum()
Out[109]:
sex
F      90993
M     110493
Name: births, dtype: int64
```

Since the dataset is split into files by year, one of the first things to do is to assemble all of the data into a single DataFrame and further add a `year` field. You can do this using `pandas.concat`. Run the following in a Jupyter cell:

```
pieces = []
for year in range(1880, 2011):
    path = f"datasets/babynames/yob{year}.txt"
    frame = pd.read_csv(path, names=["name", "sex", "births"])

    # Add a column for the year
    frame["year"] = year
    pieces.append(frame)

# Concatenate everything into a single DataFrame
names = pd.concat(pieces, ignore_index=True)
```

There are a couple things to note here. First, remember that `concat` combines the DataFrame objects by row by default. Second, you have to pass `ignore_index=True` because we're not interested in preserving the original row numbers returned from `pandas.read_csv`. So we now have a single DataFrame containing all of the names data across all years:

```
In [111]: names
Out[111]:
      name  sex  births  year
0      Mary    F     7065  1880
1      Anna    F     2604  1880
2      Emma    F     2003  1880
3  Elizabeth    F     1939  1880
4     Minnie    F     1746  1880
...
1690779   Zymaire    M      5  2010
1690780     Zyonne    M      5  2010
1690781  Zyquarius    M      5  2010
1690782      Zyran    M      5  2010
1690783     Zzyzx    M      5  2010
[1690784 rows x 4 columns]
```

With this data in hand, we can already start aggregating the data at the year and sex level using `groupby` or `pivot_table` (see [Figure 13-4](#)):

```
In [112]: total_births = names.pivot_table("births", index="year",
                                         ....,
                                         columns="sex", aggfunc=sum)

In [113]: total_births.tail()
Out[113]:
sex          F          M
year
2006  1896468  2050234
2007  1916888  2069242
2008  1883645  2032310
2009  1827643  1973359
2010  1759010  1898382

In [114]: total_births.plot(title="Total births by sex and year")
```

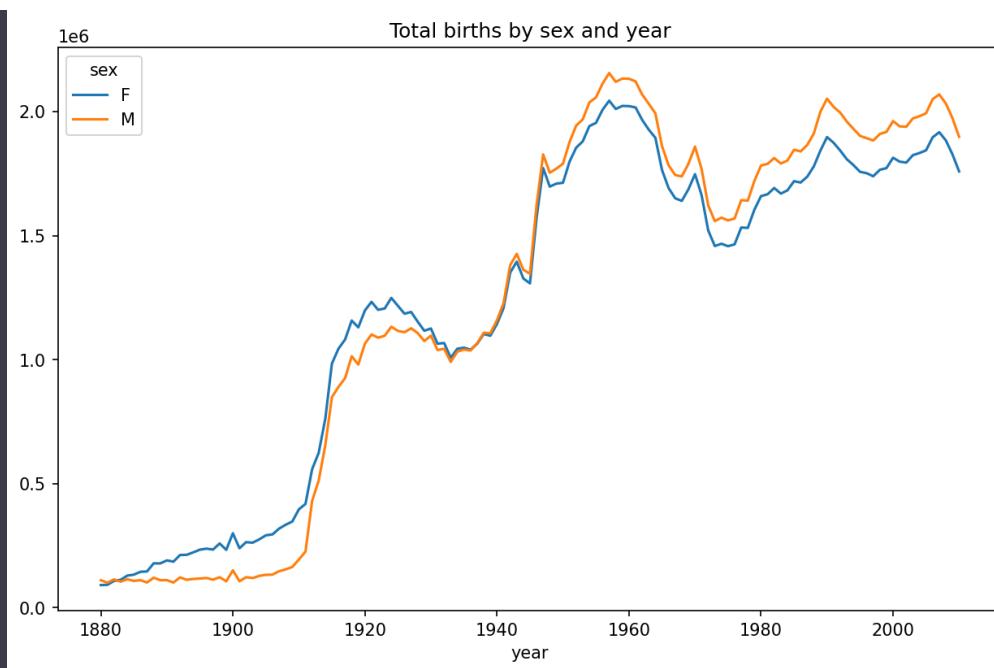


Figure 13-4. Total births by sex and year

Next, let's insert a column `prop` with the fraction of babies given each name relative to the total number of births. A `prop` value of `0.02` would indicate that 2 out of every 100 babies were given a particular name. Thus, we group the data by year and sex, then add the new column to each group:

```
def add_prop(group):
    group["prop"] = group["births"] / group["births"].sum()
    return group
names = names.groupby(["year", "sex"]).apply(add_prop)
```

The resulting complete dataset now has the following columns:

```
In [116]: names
Out[116]:
      name  sex  births   year       prop
0      Mary    F     7065  1880  0.077643
1      Anna    F     2604  1880  0.028618
2      Emma    F     2003  1880  0.022013
3  Elizabeth    F     1939  1880  0.021309
4     Minnie    F     1746  1880  0.019188
...
1690779  Zymaire    M      5  2010  0.000003
1690780   Zyonne    M      5  2010  0.000003
1690781  Zyquarius    M      5  2010  0.000003
1690782     Zyran    M      5  2010  0.000003
1690783    Zzyzx    M      5  2010  0.000003
[1690784 rows x 5 columns]
```

When performing a group operation like this, it's often valuable to do a sanity check, like verifying that the `prop` column sums to 1 within all the groups:

```
In [117]: names.groupby(["year", "sex"])["prop"].sum()
Out[117]:
year   sex
1880   F      1.0
        M      1.0
1881   F      1.0
        M      1.0
1882   F      1.0
...
2008   M      1.0
2009   F      1.0
        M      1.0
2010   F      1.0
        M      1.0
Name: prop, Length: 262, dtype: float64
```

Now that this is done, I'm going to extract a subset of the data to facilitate further analysis: the top 1,000 names for each sex/year combination. This is yet another group operation:

```
In [118]: def get_top1000(group):
    .....
    return group.sort_values("births", ascending=False)[:1000]

In [119]: grouped = names.groupby(["year", "sex"])

In [120]: top1000 = grouped.apply(get_top1000)

In [121]: top1000.head()
Out[121]:
            name  sex  births  year       prop
year sex
1880 F    0      Mary   F    7065  1880  0.077643
      1      Anna   F    2604  1880  0.028618
      2      Emma   F    2003  1880  0.022013
      3  Elizabeth   F    1939  1880  0.021309
      4     Minnie   F    1746  1880  0.019188
```

We can drop the group index since we don't need it for our analysis:

```
In [122]: top1000 = top1000.reset_index(drop=True)
```

The resulting dataset is now quite a bit smaller:

```
In [123]: top1000.head()
Out[123]:
      name  sex  births  year      prop
0     Mary    F     7065  1880  0.077643
1     Anna    F     2604  1880  0.028618
2     Emma    F     2003  1880  0.022013
3  Elizabeth    F     1939  1880  0.021309
4    Minnie    F     1746  1880  0.019188
```

We'll use this top one thousand dataset in the following investigations into the data.

Analyzing Naming Trends

With the full dataset and the top one thousand dataset in hand, we can start analyzing various naming trends of interest. First, we can split the top one thousand names into the boy and girl portions:

```
In [124]: boys = top1000[top1000["sex"] == "M"]
```

```
In [125]: girls = top1000[top1000["sex"] == "F"]
```

Simple time series, like the number of Johns or Marys for each year, can be plotted but require some manipulation to be more useful. Let's form a pivot table of the total number of births by year and name:

```
In [126]: total_births = top1000.pivot_table("births", index="year",
                                             columns="name",
                                             aggfunc=sum)
```

Now, this can be plotted for a handful of names with DataFrame's `plot` method ([Figure 13-5](#) shows the result):

```
In [127]: total_births.info()
<class 'pandas.core.frame.DataFrame'>
Int64Index: 131 entries, 1880 to 2010
Columns: 6868 entries, Aaden to Zuri
dtypes: float64(6868)
```

```
memory usage: 6.9 MB
```

```
In [128]: subset = total_births[["John", "Harry", "Mary", "Marilyn"]]

In [129]: subset.plot(subplots=True, figsize=(12, 10),
....:             title="Number of births per year")
```

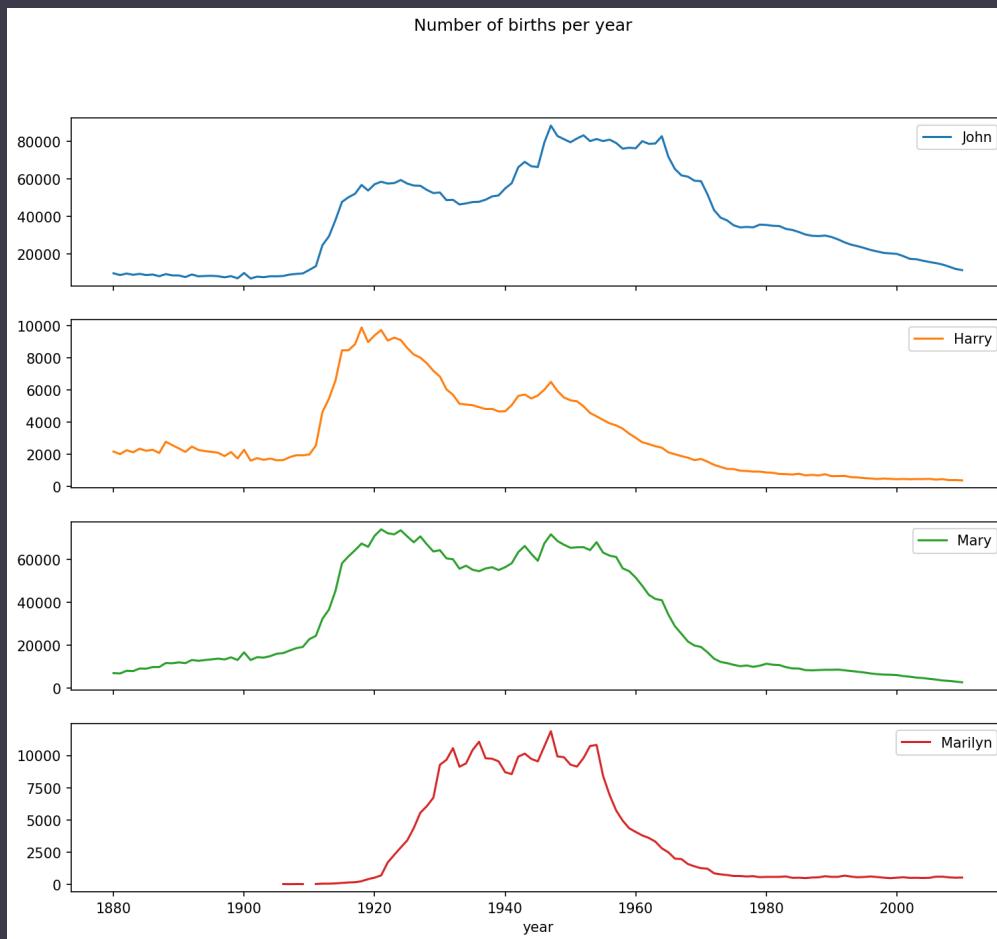


Figure 13-5. A few boy and girl names over time

On looking at this, you might conclude that these names have grown out of favor with the American population. But the story is actually more complicated than that, as will be explored in the next section.

Measuring the increase in naming diversity

One explanation for the decrease in plots is that fewer parents are choosing common names for their children. This hypothesis can be explored and confirmed in the data. One measure is the proportion of births represented by the top 1,000 most popular names, which I aggregate and plot by year and sex (Figure 13-6 shows the resulting plot):

```
In [131]: table = top1000.pivot_table("prop", index="year",
....: columns="sex", aggfunc=sum)
```

```
In [132]: table.plot(title="Sum of table1000.prop by year and sex",
.....: yticks=np.linspace(0, 1.2, 13))
```

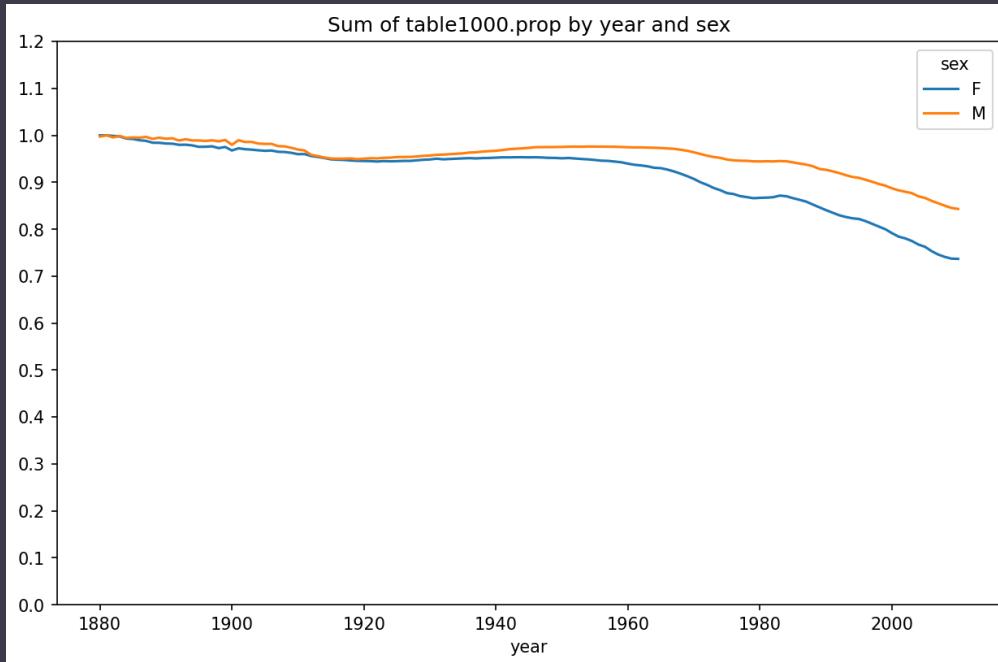


Figure 13-6. Proportion of births represented in top one thousand names by sex

You can see that, indeed, there appears to be increasing name diversity (decreasing total proportion in the top one thousand). Another interesting metric is the number of distinct names, taken in order of popularity from highest to lowest, in the top 50% of births. This number is trickier to compute. Let's consider just the boy names from 2010:

```
In [133]: df = boys[boys["year"] == 2010]
```

```
In [134]: df
```

```
Out[134]:
```

		name	sex	births	year	prop
260877		Jacob	M	21875	2010	0.011523
260878		Ethan	M	17866	2010	0.009411
260879		Michael	M	17133	2010	0.009025
260880		Jayden	M	17030	2010	0.008971
260881		William	M	16870	2010	0.008887
...
261872		Camilo	M	194	2010	0.000102
261873		Destin	M	194	2010	0.000102
261874		Jaquan	M	194	2010	0.000102
261875		Jaydan	M	194	2010	0.000102
261876		Maxton	M	193	2010	0.000102

[1000 rows x 5 columns]

After sorting `prop` in descending order, we want to know how many of the most popular names it takes to reach 50%. You could write a `for` loop to do this, but a vectorized NumPy way is more computationally efficient. Taking the cumulative sum, `cumsum`, of `prop` and then calling the method `searchsorted` returns the position in the cumulative sum at which `0.5` would need to be inserted to keep it in sorted order:

```
In [135]: prop_cumsum = df["prop"].sort_values(ascending=False).cumsum()

In [136]: prop_cumsum[:10]
Out[136]:
260877    0.011523
260878    0.020934
260879    0.029959
260880    0.038930
260881    0.047817
260882    0.056579
260883    0.065155
260884    0.073414
260885    0.081528
260886    0.089621
Name: prop, dtype: float64

In [137]: prop_cumsum.searchsorted(0.5)
Out[137]: 116
```

Since arrays are zero-indexed, adding 1 to this result gives you a result of 117. By contrast, in 1900 this number was much smaller:

```
In [138]: df = boys[boys.year == 1900]

In [139]: in1900 = df.sort_values("prop", ascending=False).prop.cumsum()

In [140]: in1900.searchsorted(0.5) + 1
Out[140]: 25
```

You can now apply this operation to each year/sex combination, `groupby` those fields, and `apply` a function returning the count for each group:

```
def get_quantile_count(group, q=0.5):
    group = group.sort_values("prop", ascending=False)
    return group.prop.cumsum().searchsorted(q) + 1
```

```
diversity = top1000.groupby(["year", "sex"]).apply(get_quantile_count)
diversity = diversity.unstack()
```

This resulting DataFrame `diversity` now has two time series, one for each sex, indexed by year. This can be inspected and plotted as before (see [Figure 13-7](#)):

```
In [143]: diversity.head()
Out[143]:
sex      F      M
year
1880    38    14
1881    38    14
1882    38    15
1883    39    15
1884    39    16
```

Figure 13-7. Plot of diversity metric by year

As you can see, girl names have always been more diverse than boy names, and they have only become more so over time. Further analysis of what exactly is driving the diversity, like the increase of alternative spellings, is left to the reader.

The “last letter” revolution

In 2007, baby name researcher Laura Wattenberg pointed out that the distribution of boy names by final letter has changed significantly over the last 100 years. To see this, we first aggregate all of the births in the full dataset by year, sex, and final letter:

Then we select three representative years spanning the history and print the first few rows:

```
In [146]: subtable = table.reindex(columns=[1910, 1960, 2010], level="year")  
  
In [147]: subtable.head()  
Out[147]:
```

sex	F			M		
year	1910	1960	2010	1910	1960	2010
last_letter						
a	108376.0	691247.0	670605.0	977.0	5204.0	28438.0
b		Nan	694.0	450.0	411.0	3912.0
c		5.0	49.0	946.0	482.0	15476.0
d		6750.0	3729.0	2607.0	22111.0	262112.0
e		133569.0	435013.0	313833.0	28655.0	178823.0

Next, normalize the table by total births to compute a new table containing the proportion of total births for each sex ending in each letter:

```
In [148]: subtable.sum()  
Out[148]:
```

sex	year	
F	1910	396416.0
	1960	2022062.0
	2010	1759010.0
M	1910	194198.0
	1960	2132588.0
	2010	1898382.0

```
dtype: float64  
  
In [149]: letter_prop = subtable / subtable.sum()  
  
In [150]: letter_prop  
Out[150]:
```

sex	F			M		
year	1910	1960	2010	1910	1960	2010
last_letter						
a	0.273390	0.341853	0.381240	0.005031	0.002440	0.014980
b		Nan	0.000343	0.000256	0.002116	0.001834
c		0.000013	0.000024	0.000538	0.002482	0.007257
d		0.017028	0.001844	0.001482	0.113858	0.122908
e		0.336941	0.215133	0.178415	0.147556	0.083853
...
v		Nan	0.000060	0.000117	0.000113	0.000037
w		0.000020	0.000031	0.001182	0.006329	0.007711
x		0.000015	0.000037	0.000727	0.003965	0.001851

```

y          0.110972  0.152569  0.116828  0.077349  0.160987  0.058168
z          0.002439  0.000659  0.000704  0.000170  0.000184  0.001831
[26 rows x 6 columns]

```

With the letter proportions now in hand, we can make bar plots for each sex, broken down by year (see [Figure 13-8](#)):

```

import matplotlib.pyplot as plt

fig, axes = plt.subplots(2, 1, figsize=(10, 8))
letter_prop["M"].plot(kind="bar", rot=0, ax=axes[0], title="Male")
letter_prop["F"].plot(kind="bar", rot=0, ax=axes[1], title="Female",
                      legend=False)

```

Figure 13-8. Proportion of boy and girl names ending in each letter

As you can see, boy names ending in *n* have experienced significant growth since the 1960s. Going back to the full table created before, I again normalize by year and sex and select a subset of letters for the boy names, finally transposing to make each column a time series:

```

In [153]: letter_prop = table / table.sum()

In [154]: dny_ts = letter_prop.loc[["d", "n", "y"], "M"].T

In [155]: dny_ts.head()
Out[155]:
last_letter      d      n      y
year
1880      0.083055  0.153213  0.075760
1881      0.083247  0.153214  0.077451
1882      0.085340  0.149560  0.077537
1883      0.084066  0.151646  0.079144
1884      0.086120  0.149915  0.080405

```

With this DataFrame of time series in hand, I can make a plot of the trends over time again with its `plot` method (see [Figure 13-9](#)):

```
In [158]: dny_ts.plot()
```

Figure 13-9. Proportion of boys born with names ending in d/n/y over time

Boy names that became girl names (and vice versa)

Another fun trend is looking at names that were more popular with one gender earlier in the sample but have become preferred as a name for the other gender over time. One example is the name Lesley or Leslie. Going back to the `top1000` DataFrame, I compute a list of names occurring in the dataset starting with “Lesl”:

```
In [159]: all_names = pd.Series(top1000["name"].unique())

In [160]: lesley_like = all_names[all_names.str.contains("Lesl")]

In [161]: lesley_like
Out[161]:
632      Leslie
2294     Lesley
4262    Leslee
4728     Lesli
6103     Lesly
dtype: object
```

From there, we can filter down to just those names and sum births grouped by name to see the relative frequencies:

```
In [162]: filtered = top1000[top1000["name"].isin(lesley_like)]

In [163]: filtered.groupby("name")["births"].sum()
Out[163]:
name
Leslee      1082
Lesley     35022
Lesli       929
Leslie    370429
Lesly      10067
Name: births, dtype: int64
```

Next, let’s aggregate by sex and year, and normalize within year:

```
In [164]: table = filtered.pivot_table("births", index="year",
                                         ....,
                                         columns="sex", aggfunc="sum")

In [165]: table = table.div(table.sum(axis="columns"), axis="index")

In [166]: table.tail()
```

```
Out[166]:  
sex      F      M  
year  
2006   1.0  NaN  
2007   1.0  NaN  
2008   1.0  NaN  
2009   1.0  NaN  
2010   1.0  NaN
```

Lastly, it's now possible to make a plot of the breakdown by sex over time (see [Figure 13-10](#)):

```
In [168]: table.plot(style={"M": "k-", "F": "k--"})
```

Figure 13-10. Proportion of male/female Lesley-like names over time

13.4 USDA Food Database

The US Department of Agriculture (USDA) makes available a database of food nutrient information. Programmer Ashley Williams created a version of this database in JSON format. The records look like this:

```
{  
    "id": 21441,  
    "description": "KENTUCKY FRIED CHICKEN, Fried Chicken, EXTRA CRISPY,  
    Wing, meat and skin with breading",  
    "tags": ["KFC"],  
    "manufacturer": "Kentucky Fried Chicken",  
    "group": "Fast Foods",  
    "portions": [  
        {  
            "amount": 1,  
            "unit": "wing, with skin",  
            "grams": 68.0  
        },  
        ...  
    ],  
    "nutrients": [  
        {  
            "value": 20.8,  
            "units": "g",  
            "description": "Protein",  
            "name": "Protein"  
        }  
    ]  
}
```

```
        "group": "Composition"
    },
    ...
]
```

Each food has a number of identifying attributes along with two lists of nutrients and portion sizes. Data in this form is not particularly amenable to analysis, so we need to do some work to wrangle the data into a better form.

You can load this file into Python with any JSON library of your choosing. I'll use the built-in Python `json` module:

```
In [169]: import json

In [170]: db = json.load(open("datasets/usda_food/database.json"))

In [171]: len(db)
Out[171]: 6636
```

Each entry in `db` is a dictionary containing all the data for a single food. The `"nutrients"` field is a list of dictionaries, one for each nutrient:

```
In [172]: db[0].keys()
Out[172]: dict_keys(['id', 'description', 'tags', 'manufacturer', 'group', 'ons', 'nutrients'])

In [173]: db[0]["nutrients"][0]
Out[173]:
{'value': 25.18,
 'units': 'g',
 'description': 'Protein',
 'group': 'Composition'}

In [174]: nutrients = pd.DataFrame(db[0]["nutrients"])

In [175]: nutrients.head(7)
Out[175]:
   value  units           description      group
0   25.18     g             Protein  Composition
1   29.20     g  Total lipid (fat)  Composition
2   3.06     g  Carbohydrate, by difference  Composition
3   3.28     g                 Ash          Other
```

4	376.00	kcal	Energy	Energy
5	39.28	g	Water	Composition
6	1573.00	kJ	Energy	Energy

When converting a list of dictionaries to a DataFrame, we can specify a list of fields to extract. We'll take the food names, group, ID, and manufacturer:

```
In [176]: info_keys = ["description", "group", "id", "manufacturer"]

In [177]: info = pd.DataFrame(db, columns=info_keys)

In [178]: info.head()
Out[178]:
      description          group    id
0   Cheese, caraway  Dairy and Egg Products  1008
1   Cheese, cheddar  Dairy and Egg Products  1009
2   Cheese, edam    Dairy and Egg Products  1018
3   Cheese, feta    Dairy and Egg Products  1019
4 Cheese, mozzarella, part skim milk  Dairy and Egg Products  1028

      manufacturer
0
1
2
3
4

In [179]: info.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 6636 entries, 0 to 6635
Data columns (total 4 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   description      6636 non-null   object 
 1   group            6636 non-null   object 
 2   id               6636 non-null   int64  
 3   manufacturer     5195 non-null   object 
dtypes: int64(1), object(3)
memory usage: 207.5+ KB
```

From the output of `info.info()`, we can see that there is missing data in the `manufacturer` column.

You can see the distribution of food groups with `value_counts`:

```
In [180]: pd.value_counts(info["group"])[10]
Out[180]:
Vegetables and Vegetable Products    812
Beef Products                      618
Baked Products                     496
Breakfast Cereals                  403
Legumes and Legume Products        365
Fast Foods                         365
Lamb, Veal, and Game Products     345
Sweets                            341
Fruits and Fruit Juices          328
Pork Products                      328
Name: group, dtype: int64
```

Now, to do some analysis on all of the nutrient data, it's easiest to assemble the nutrients for each food into a single large table. To do so, we need to take several steps. First, I'll convert each list of food nutrients to a DataFrame, add a column for the food `id`, and append the DataFrame to a list. Then, these can be concatenated with `concat`. Run the following code in a Jupyter cell:

```
nutrients = []

for rec in db:
    fnuts = pd.DataFrame(rec["nutrients"])
    fnuts["id"] = rec["id"]
    nutrients.append(fnuts)

nutrients = pd.concat(nutrients, ignore_index=True)
```

If all goes well, `nutrients` should look like this:

```
In [182]: nutrients
Out[182]:
   value units           description      group
0   25.180   g             Protein Composition 100
1   29.200   g       Total lipid (fat) Composition 100
2   3.060   g  Carbohydrate, by difference Composition 100
3   3.280   g                 Ash            Other 100
4  376.000   kcal            Energy            Energy 100
... ...
389350  0.000   mcg  Vitamin B-12, added Vitamins 43540
389351  0.000   mg        Cholesterol            Other 43540
389352  0.072   g  Fatty acids, total saturated            Other 43540
389353  0.028   g Fatty acids, total monounsaturated            Other 43540
```

```
389354    0.041      g Fatty acids, total polyunsaturated  
[389355 rows x 5 columns]
```

Other 4354

I noticed that there are duplicates in this DataFrame, so it makes things easier to drop them:

```
In [183]: nutrients.duplicated().sum() # number of duplicates  
Out[183]: 14179  
  
In [184]: nutrients = nutrients.drop_duplicates()
```

Since "group" and "description" are in both DataFrame objects, we can rename for clarity:

```
In [185]: col_mapping = {"description" : "food",  
.....: "group" : "fgroup"}  
  
In [186]: info = info.rename(columns=col_mapping, copy=False)  
  
In [187]: info.info()  
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 6636 entries, 0 to 6635  
Data columns (total 4 columns):  
 #   Column           Non-Null Count  Dtype     
---  --  
 0   food             6636 non-null    object    
 1   fgroup           6636 non-null    object    
 2   id               6636 non-null    int64     
 3   manufacturer     5195 non-null    object    
dtypes: int64(1), object(3)  
memory usage: 207.5+ KB  
  
In [188]: col_mapping = {"description" : "nutrient",  
.....: "group" : "nutgroup"}  
  
In [189]: nutrients = nutrients.rename(columns=col_mapping, copy=False)
```

```
In [190]: nutrients  
Out[190]:  
      value units          nutrient      nutgroup  
0      25.180   g            Protein    Composition  
1      29.200   g            Total lipid (fat) Composition  
2      3.060   g Carbohydrate, by difference Composition  
3      3.280   g                  Ash        Other  
4     376.000   kcal                 Energy    Energy  
..       .. .
```

```

389350    0.000    mcg          Vitamin B-12, added
389351    0.000    mg           Cholesterol
389352    0.072    g            Fatty acids, total saturated
389353    0.028    g            Fatty acids, total monounsaturated
389354    0.041    g            Fatty acids, total polyunsaturated
[375176 rows x 5 columns]

```

Vitamins	4354
Other	4354

With all of this done, we're ready to merge `info` with `nutrients`:

```

In [191]: ndata = pd.merge(nutrients, info, on="id")

In [192]: ndata.info()
<class 'pandas.core.frame.DataFrame'>
Int64Index: 375176 entries, 0 to 375175
Data columns (total 8 columns):
 #   Column        Non-Null Count  Dtype  
--- 
 0   value         375176 non-null   float64
 1   units         375176 non-null   object 
 2   nutrient      375176 non-null   object 
 3   nutgroup      375176 non-null   object 
 4   id             375176 non-null   int64  
 5   food           375176 non-null   object 
 6   fgroup         375176 non-null   object 
 7   manufacturer   293054 non-null   object 
dtypes: float64(1), int64(1), object(6)
memory usage: 25.8+ MB

In [193]: ndata.iloc[30000]
Out[193]:
value                      0.04
units                       g
nutrient                   Glycine
nutgroup                  Amino Acids
id                         6158
food           Soup, tomato bisque, canned, condensed
fgroup          Soups, Sauces, and Gravies
manufacturer
Name: 30000, dtype: object

```

We could now make a plot of median values by food group and nutrient type (see [Figure 13-11](#)):

```

In [195]: result = ndata.groupby(["nutrient", "fgroup"])["value"].quantile(0.5)

In [196]: result["Zinc, Zn"].sort_values().plot(kind="barh")

```

Figure 13-11. Median zinc values by food group

Using the `idxmax` or `argmax` Series methods, you can find which food is most dense in each nutrient. Run the following in a Jupyter cell:

```
by_nutrient = ndata.groupby(["nutgroup", "nutrient"])

def get_maximum(x):
    return x.loc[x.value.idxmax()]

max_foods = by_nutrient.apply(get_maximum)[["value", "food"]]

# make the food a little smaller
max_foods["food"] = max_foods["food"].str[:50]
```

The resulting DataFrame is a bit too large to display in the book; here is only the `"Amino Acids"` nutrient group:

```
In [198]: max_foods.loc["Amino Acids"]["food"]
Out[198]:
nutrient
Alanine           Gelatins, dry powder, unsweetened
Arginine          Seeds, sesame flour, low-fat
Aspartic acid     Soy protein isolate
Cystine           Seeds, cottonseed flour, low fat (glandless)
Glutamic acid    Soy protein isolate
Glycine           Gelatins, dry powder, unsweetened
Histidine         Whale, beluga, meat, dried (Alaska Native)
Hydroxyproline   KENTUCKY FRIED CHICKEN, Fried Chicken, ORIGINAL RE
Isoleucine        Soy protein isolate, PROTEIN TECHNOLOGIES INTERNAT
Leucine           Soy protein isolate, PROTEIN TECHNOLOGIES INTERNAT
Lysine            Seal, bearded (Oogruk), meat, dried (Alaska Native)
Methionine        Fish, cod, Atlantic, dried and salted
Phenylalanine    Soy protein isolate, PROTEIN TECHNOLOGIES INTERNAT
Proline           Gelatins, dry powder, unsweetened
Serine            Soy protein isolate, PROTEIN TECHNOLOGIES INTERNAT
Threonine         Soy protein isolate, PROTEIN TECHNOLOGIES INTERNAT
Tryptophan        Sea lion, Steller, meat with fat (Alaska Native)
Tyrosine          Soy protein isolate, PROTEIN TECHNOLOGIES INTERNAT
Valine            Soy protein isolate, PROTEIN TECHNOLOGIES INTERNAT
Name: food, dtype: object
```

13.5 2012 Federal Election Commission Database

The US Federal Election Commission (FEC) publishes data on contributions to political campaigns. This includes contributor names, occupation and employer, address, and contribution amount. The contribution data from the 2012 US presidential election was available as a single 150-megabyte CSV file *P00000001-ALL.csv* (see the book's data repository), which can be loaded with `pandas.read_csv`:

```
In [199]: fec = pd.read_csv("datasets/fec/P00000001-ALL.csv", low_memory=False)

In [200]: fec.info()
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1001731 entries, 0 to 1001730
Data columns (total 16 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   cmte_id          1001731 non-null   object 
 1   cand_id          1001731 non-null   object 
 2   cand_nm          1001731 non-null   object 
 3   contbr_nm        1001731 non-null   object 
 4   contbr_city      1001712 non-null   object 
 5   contbr_st        1001727 non-null   object 
 6   contbr_zip       1001620 non-null   object 
 7   contbr_employer   988002 non-null   object 
 8   contbr_occupation 993301 non-null   object 
 9   contb_receipt_amt 1001731 non-null   float64
 10  contb_receipt_dt 1001731 non-null   object 
 11  receipt_desc     14166 non-null    object 
 12  memo_cd          92482 non-null   object 
 13  memo_text         97770 non-null   object 
 14  form_tp          1001731 non-null   object 
 15  file_num          1001731 non-null   int64  
dtypes: float64(1), int64(1), object(14)
memory usage: 122.3+ MB
```

NOTE

Several people asked me to update the dataset from the 2012 election to the 2016 or 2020 elections. Unfortunately, the more recent datasets provided by the FEC have become larger and more complex, and I decided that working with them here would be a distraction from the analysis techniques that I wanted to illustrate.

A sample record in the DataFrame looks like this:

```
In [201]: fec.iloc[123456]
Out[201]:
cmte_id           C00431445
cand_id           P80003338
cand_nm          Obama, Barack
contbr_nm         ELLMAN, IRA
contbr_city       TEMPE
contbr_st          AZ
contbr_zip        852816719
contbr_employer   ARIZONA STATE UNIVERSITY
contbr_occupation PROFESSOR
contb_receipt_amt    50.0
contb_receipt_dt  01-DEC-11
receipt_desc      NaN
memo_cd           NaN
memo_text          NaN
form_tp            SA17A
file_num           772372
Name: 123456, dtype: object
```

You may think of some ways to start slicing and dicing this data to extract informative statistics about donors and patterns in the campaign contributions. I'll show you a number of different analyses that apply the techniques in this book.

You can see that there are no political party affiliations in the data, so this would be useful to add. You can get a list of all the unique political candidates using `unique`:

```
In [202]: unique_cands = fec["cand_nm"].unique()

In [203]: unique_cands
Out[203]:
array(['Bachmann, Michelle', 'Romney, Mitt', 'Obama, Barack'],
```

```
"Roemer, Charles E. 'Buddy' III", "Pawlenty, Timothy",  
"Johnson, Gary Earl", "Paul, Ron", "Santorum, Rick",  
"Cain, Herman", "Gingrich, Newt", "McCotter, Thaddeus G",  
"Huntsman, Jon", "Perry, Rick"], dtype=object)
```

```
In [204]: unique_cands[2]  
Out[204]: 'Obama, Barack'
```

One way to indicate party affiliation is using a dictionary:¹

```
parties = {"Bachmann, Michelle": "Republican",  
           "Cain, Herman": "Republican",  
           "Gingrich, Newt": "Republican",  
           "Huntsman, Jon": "Republican",  
           "Johnson, Gary Earl": "Republican",  
           "McCotter, Thaddeus G": "Republican",  
           "Obama, Barack": "Democrat",  
           "Paul, Ron": "Republican",  
           "Pawlenty, Timothy": "Republican",  
           "Perry, Rick": "Republican",  
           "Roemer, Charles E. 'Buddy' III": "Republican",  
           "Romney, Mitt": "Republican",  
           "Santorum, Rick": "Republican"}
```

Now, using this mapping and the `map` method on Series objects, you can compute an array of political parties from the candidate names:

```
In [206]: fec["cand_nm"][123456:123461]  
Out[206]:  
123456    Obama, Barack  
123457    Obama, Barack  
123458    Obama, Barack  
123459    Obama, Barack  
123460    Obama, Barack  
Name: cand_nm, dtype: object  
  
In [207]: fec["cand_nm"][123456:123461].map(parties)  
Out[207]:  
123456    Democrat  
123457    Democrat  
123458    Democrat  
123459    Democrat  
123460    Democrat  
Name: cand_nm, dtype: object  
  
# Add it as a column
```

```
In [208]: fec["party"] = fec["cand_nm"].map(parties)

In [209]: fec["party"].value_counts()
Out[209]:
Democrat      593746
Republican    407985
Name: party, dtype: int64
```

A couple of data preparation points. First, this data includes both contributions and refunds (negative contribution amount):

```
In [210]: (fec["contb_receipt_amt"] > 0).value_counts()
Out[210]:
True      991475
False     10256
Name: contb_receipt_amt, dtype: int64
```

To simplify the analysis, I'll restrict the dataset to positive contributions:

```
In [211]: fec = fec[fec["contb_receipt_amt"] > 0]
```

Since Barack Obama and Mitt Romney were the main two candidates, I'll also prepare a subset that just has contributions to their campaigns:

```
In [212]: fec_mrbo = fec[fec["cand_nm"].isin(["Obama, Barack", "Romney, Mitt
```

Donation Statistics by Occupation and Employer

Donations by occupation is another oft-studied statistic. For example, attorneys tend to donate more money to Democrats, while business executives tend to donate more to Republicans. You have no reason to believe me; you can see for yourself in the data. First, the total number of donations by occupation can be computed with `value_counts`:

```
In [213]: fec["contbr_occupation"].value_counts()[:10]
Out[213]:
RETIRED                  233990
INFORMATION REQUESTED    35107
ATTORNEY                 34286
HOMEMAKER                29931
PHYSICIAN                 23432
INFORMATION REQUESTED PER BEST EFFORTS  21138
```

ENGINEER	14334
TEACHER	13990
CONSULTANT	13273
PROFESSOR	12555
Name: contbr_occupation, dtype: int64	

You will notice by looking at the occupations that many refer to the same basic job type, or there are several variants of the same thing. The following code snippet illustrates a technique for cleaning up a few of them by mapping from one occupation to another; note the “trick” of using `dict.get` to allow occupations with no mapping to “pass through”:

```
occ_mapping = {
    "INFORMATION REQUESTED PER BEST EFFORTS" : "NOT PROVIDED",
    "INFORMATION REQUESTED" : "NOT PROVIDED",
    "INFORMATION REQUESTED (BEST EFFORTS)" : "NOT PROVIDED",
    "C.E.O." : "CEO"
}

def get_occ(x):
    # If no mapping provided, return x
    return occ_mapping.get(x, x)

fec["contbr_occupation"] = fec["contbr_occupation"].map(get_occ)
```

I'll also do the same thing for employers:

```
emp_mapping = {
    "INFORMATION REQUESTED PER BEST EFFORTS" : "NOT PROVIDED",
    "INFORMATION REQUESTED" : "NOT PROVIDED",
    "SELF" : "SELF-EMPLOYED",
    "SELF EMPLOYED" : "SELF-EMPLOYED",
}

def get_emp(x):
    # If no mapping provided, return x
    return emp_mapping.get(x, x)

fec["contbr_employer"] = fec["contbr_employer"].map(f)
```

Now, you can use `pivot_table` to aggregate the data by party and occupation, then filter down to the subset that donated at least \$2 million overall:

```
In [216]: by_occupation = fec.pivot_table("contb_receipt_amt",
                                         ....,
                                         ....,
                                         index="contbr_occupation",
                                         columns="party",
                                         aggfunc="sum")

In [217]: over_2mm = by_occupation[by_occupation.sum(axis="columns") > 200000]

In [218]: over_2mm
Out[218]:
party           Democrat   Republican
contbr_occupation
ATTORNEY        11141982.97    7477194.43
CEO              2074974.79    4211040.52
CONSULTANT       2459912.71    2544725.45
ENGINEER         951525.55    1818373.70
EXECUTIVE        1355161.05    4138850.09
HOMEMAKER        4248875.80    13634275.78
INVESTOR          884133.00    2431768.92
LAWYER            3160478.87    391224.32
MANAGER           762883.22    1444532.37
NOT PROVIDED     4866973.96    20565473.01
OWNER             1001567.36    2408286.92
PHYSICIAN         3735124.94    3594320.24
PRESIDENT         1878509.95    4720923.76
PROFESSOR         2165071.08    296702.73
REAL ESTATE        528902.09    1625902.25
RETIRED            25305116.38   23561244.49
SELF-EMPLOYED      672393.40    1640252.54
```

It can be easier to look at this data graphically as a bar plot ("barh" means horizontal bar plot; see [Figure 13-12](#)):

```
In [220]: over_2mm.plot(kind="barh")
```

Figure 13-12. Total donations by party for top occupations

You might be interested in the top donor occupations or top companies that donated to Obama and Romney. To do this, you can group by candidate name and use a variant of the `top` method from earlier in the chapter:

```
def get_top_amounts(group, key, n=5):
    totals = group.groupby(key)[["contb_receipt_amt"]].sum()
    return totals.nlargest(n)
```

Then aggregate by occupation and employer:

```
In [222]: grouped = fec_mrbo.groupby("cand_nm")

In [223]: grouped.apply(get_top_amounts, "contbr_occupation", n=7)
Out[223]:
cand_nm      contbr_occupation
Obama, Barack    RETIRED           25305116.38
                  ATTORNEY         11141982.97
                  INFORMATION REQUESTED 4866973.96
                  HOMEMAKER        4248875.80
                  PHYSICIAN         3735124.94
                  LAWYER            3160478.87
                  CONSULTANT        2459912.71
Romney, Mitt     RETIRED           11508473.59
                  INFORMATION REQUESTED PER BEST EFFORTS 11396894.84
                  HOMEMAKER        8147446.22
                  ATTORNEY          5364718.82
                  PRESIDENT          2491244.89
                  EXECUTIVE          2300947.03
                  C.E.O.             1968386.11
Name: contb_receipt_amt, dtype: float64

In [224]: grouped.apply(get_top_amounts, "contbr_employer", n=10)
Out[224]:
cand_nm      contbr_employer
Obama, Barack    RETIRED           22694358.85
                  SELF-EMPLOYED      17080985.96
                  NOT EMPLOYED       8586308.70
                  INFORMATION REQUESTED 5053480.37
                  HOMEMAKER          2605408.54
                  SELF               1076531.20
                  SELF EMPLOYED      469290.00
                  STUDENT            318831.45
                  VOLUNTEER          257104.00
                  MICROSOFT          215585.36
Romney, Mitt     INFORMATION REQUESTED PER BEST EFFORTS 12059527.24
                  RETIRED            11506225.71
                  HOMEMAKER          8147196.22
                  SELF-EMPLOYED      7409860.98
                  STUDENT            496490.94
                  CREDIT SUISSE        281150.00
                  MORGAN STANLEY      267266.00
                  GOLDMAN SACH & CO.   238250.00
                  BARCLAYS CAPITAL     162750.00
                  H.I.G. CAPITAL        139500.00
Name: contb_receipt_amt, dtype: float64
```

Bucketing Donation Amounts

A useful way to analyze this data is to use the `cut` function to discretize the contributor amounts into buckets by contribution size:

```
In [225]: bins = np.array([0, 1, 10, 100, 1000, 10000,
...:                      100_000, 1_000_000, 10_000_000])
In [226]: labels = pd.cut(fec_mrbo["contb_receipt_amt"], bins)
In [227]: labels
Out[227]:
411      (10, 100]
412      (100, 1000]
413      (100, 1000]
414      (10, 100]
415      (10, 100]
...
701381    (10, 100]
701382    (100, 1000]
701383    (1, 10]
701384    (10, 100]
701385    (100, 1000]
Name: contb_receipt_amt, Length: 694282, dtype: category
Categories (8, interval[int64, right]): [(0, 1] < (1, 10] < (10, 100] < (100, 1000] <
(1000, 10000] < (10000, 100000] < (100000, 1000000] <
(1000000, 10000000]
```

We can then group the data for Obama and Romney by name and bin label to get a histogram by donation size:

```
In [228]: grouped = fec_mrbo.groupby(["cand_nm", "labels"])
In [229]: grouped.size().unstack(level=0)
Out[229]:
cand_nm          Obama, Barack  Romney, Mitt
contb_receipt_amt
(0, 1]                493        77
(1, 10]               40070      3681
(10, 100]              372280     31853
(100, 1000]             153991     43357
(1000, 10000]            22284     26186
(10000, 100000]                2           1
```

(100000, 1000000]	3	0
(1000000, 10000000]	4	0

This data shows that Obama received a significantly larger number of small donations than Romney. You can also sum the contribution amounts and normalize within buckets to visualize the percentage of total donations of each size by candidate ([Figure 13-13](#) shows the resulting plot):

```
In [231]: bucket_sums = grouped["contb_receipt_amt"].sum().unstack(level=0)

In [232]: normed_sums = bucket_sums.div(bucket_sums.sum(axis="columns"),
.....:                               axis="index")

In [233]: normed_sums
Out[233]:
cand_nm          Obama, Barack  Romney, Mitt
contb_receipt_amt
(0, 1]           0.805182    0.194818
(1, 10]          0.918767    0.081233
(10, 100]         0.910769    0.089231
(100, 1000]       0.710176    0.289824
(1000, 10000]     0.447326    0.552674
(10000, 100000]   0.823120    0.176880
(100000, 1000000] 1.000000    0.000000
(1000000, 10000000] 1.000000    0.000000

In [234]: normed_sums[:-2].plot(kind="barh")
```

Figure 13-13. Percentage of total donations received by candidates for each donation size

I excluded the two largest bins, as these are not donations by individuals.

This analysis can be refined and improved in many ways. For example, you could aggregate donations by donor name and zip code to adjust for donors who gave many small amounts versus one or more large donations. I encourage you to explore the dataset yourself.

Donation Statistics by State

We can start by aggregating the data by candidate and state:

```
In [235]: grouped = fec_mrbo.groupby(["cand_nm", "contbr_st"])

In [236]: totals = grouped["contb_receipt_amt"].sum().unstack(level=0).fillna(0)

In [237]: totals = totals[totals.sum(axis="columns") > 1000000]

In [238]: totals.head(10)
Out[238]:
cand_nm      Obama, Barack    Romney, Mitt
contbr_st
AK           281840.15     86204.24
AL           543123.48    527303.51
AR           359247.28    105556.00
AZ           1506476.98   1888436.23
CA           23824984.24  11237636.60
CO           2132429.49   1506714.12
CT           2068291.26   3499475.45
DC           4373538.80   1025137.50
DE           336669.14    82712.00
FL           7318178.58   8338458.81
```

If you divide each row by the total contribution amount, you get the relative percentage of total donations by state for each candidate:

```
In [239]: percent = totals.div(totals.sum(axis="columns")), axis="index")

In [240]: percent.head(10)
Out[240]:
cand_nm      Obama, Barack    Romney, Mitt
contbr_st
AK           0.765778     0.234222
AL           0.507390     0.492610
AR           0.772902     0.227098
AZ           0.443745     0.556255
CA           0.679498     0.320502
CO           0.585970     0.414030
CT           0.371476     0.628524
DC           0.810113     0.189887
DE           0.802776     0.197224
FL           0.467417     0.532583
```

13.6 Conclusion

We've reached the end of this book. I have included some additional content you may find useful in the appendixes.

In the 10 years since the first edition of this book was published, Python has become a popular and widespread language for data analysis. The programming skills you have developed here will stay relevant for a long time into the future. I hope the programming tools and libraries we've explored will serve you well.

- 1 This makes the simplifying assumption that Gary Johnson is a Republican even though he later became the Libertarian party candidate.

[Support](#) | [Sign Out](#)

©2022 O'REILLY MEDIA, INC. [TERMS OF SERVICE](#) | [PRIVACY POLICY](#)